

Logs cost the most and get read the least

Metrics, logs and traces answer different questions and have different cost curves. One line item hides all three.

PRESENTER · TEAM

Three signals answer three different questions

	Metrics	Logs	Traces
ANSWERS	Is it happening, and how much	What the code actually did	Where the time went, hop by hop
BLIND TO	Which request. Which user.	Where the time went	What the code was thinking
GROWS WITH	Label cardinality, not traffic	Traffic, replicas and days kept	Sampled share, span width

Volume grows with traffic. Readership does not.

WHAT MULTIPLIES

Requests per second

- × Log lines emitted per request
- × Replicas emitting the same line
- × Days retained, indexed, replicated

WHAT DOES NOT

Incidents, one at a time

Engineers who open logs at all

The window they read: minutes

The three queries they know

Cost is a **product**. Readership is a **constant**. That gap is the whole argument.

Sampling and retention are bets, not savings

THE TRADE	WHAT IT KEEPS	WHAT IT SPENDS
Sampling	The shape of normal traffic	The one request that mattered
Retention	What is happening right now	Whether it happened before

You find out which bet you took at 03:00, from a query that returns nothing.
Nobody called it a bet when it was placed.

You cannot argue with one line item

Observability

One line. One owner. One available question: keep it or cut it.

BY SIGNAL

Metrics, logs and traces
have different curves.
Price them apart.

BY VERBOSITY

Debug level in
production is a choice.
Someone has to make it.

BY WINDOW

Hot days, warm weeks,
cold archive — three
prices, not one.

One line item can only be approved or refused. Three can be negotiated.

What we need to decide

Three decisions. Each one makes part
of the bill arguable.

1 What gets sampled

Which signals, at what share, changed by whom

2 What gets dropped

Which lines never leave the process at all

3 Who owns retention

One named owner for each window we pay for

 None of the three is a tooling decision.